

How we use Go to build APIs

PubNative best practices and patterns

Kostiantyn Stepaniuk

PubNative GmbH

1/17 About PubNative

- Adtech company which connects publishers and advertisers through its technology
- First line of the API code was written in **December 2013** in **Ruby**
- In **November 2014** we migrated to **Go**
- API and supportive utilities have **200K** lines of **Go** code
- We process **12M req/min**
- Response time below **100ms**
- We are connected to **197** partners with who we are doing real-time auction
- At a peak, we send **$12M * 197 = 2.364B$** outbound req/min

2/17 What challenges we have

- **Competitive:** cost-efficient and high-performing
- **Resilient:** survives the outage of any used service, DB, queue, cache, etc.
- **Quick TTM:** 1-2 products/year
- **Productivity:** 3043 PRs closed since 2013 or 2 PRs/day
- **Efficiency:** product/tech task split 85/15, developing by 4 people (of 14 engineers)

3/17 How we do it

Make it simple

Make it boring

Focus on what matters

4/17 Learn techniques



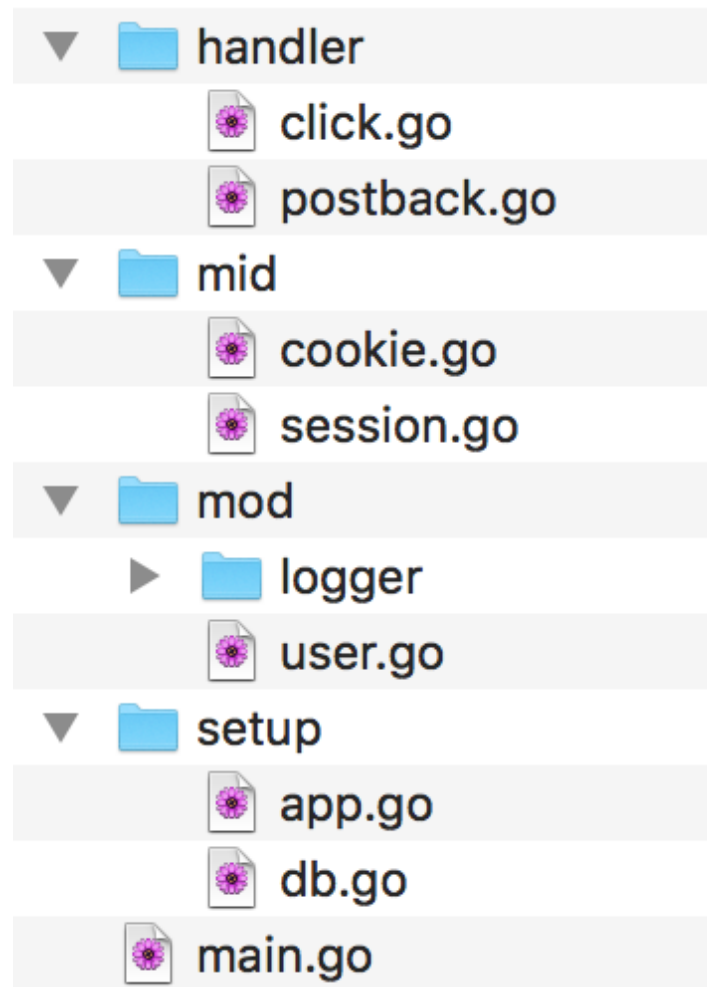
Painted in 1897 by 15-year-old boy

5/17 And then improve



Painted in 1955 by our 74-year-old boy :)

6/17 Use the same structure



7/17 Full control over HTTP protocol

```
http.HandleFunc("/handle/b", func(http.ResponseWriter, *http.Request) {})  
http.HandleFunc("/handle/a", func(http.ResponseWriter, *http.Request) {})  
  
http.HandleFunc("/basic_auth", func(w http.ResponseWriter, r *http.Request) {  
    usr, pwd, _ := r.BasicAuth()  
    if usr != "user" || pwd != "password" {  
        w.Header().Set("WWW-Authenticate", `Basic realm="Application"`)  
        w.WriteHeader(401)  
    }  
})
```

- No powerful URL routers like [gorilla/mux](https://github.com/gorilla/mux) because of performance penalty and decreasing clarity. Go provides all the necessary wrappers
- Don't carry parameters in the path, use a query string or a request payload
- We need to know what happens under the hood to optimize

8/17 Use explicit types, avoid interfaces

```
type App struct {  
    Redis *redis.Client      // explicit name  
    HB    *honeybadger.Client //commonly used abbreviation  
}
```

- We solve specific problems with concrete tools
- Clear understanding what is used and how
- Renaming can be done in one command in your favorite editor

```
type ObjectA struct{} // Do we really need to have one interface?  
type ObjectB struct{} // Or we can keep them separately  
func (o ObjectA) Run(task int) bool {}  
func (o ObjectB) Do(task int) bool {}
```

- Got a challenge, try to solve it in the best way possible and then think of generalizing, not another way around

9/17 One way of doing things

```
type A struct {  
    B B  
}  
type B struct {  
    C C  
}  
type C struct {  
    D string  
}  
  
a := A{  
a.B.C.D // access D through chain  
  
a.GetD() // NO helper methods  
  
a.C.D // NO aliases
```

It also applies how to create pools, integrates 3rd parties or writes tests, etc.

10/17 Own pool of workers instead of what libraries provide

```
func main() {  
    pool := 10  
    ch := make(chan string, pool+10) // how many messages keep in a memory before taking an action  
    for i := 0; i < pool; i++ {  
        conn, err := redis.Dial("tcp", ":6379") // connection per worker  
        _ = err // handle error  
        go func() {  
            conn.Do("HSET", "hash", <-ch, 1)  
        }()  
    }  
}
```

- Know how many connections are used
- No hidden bottlenecks, e.g., locks are used to manage the pool
- Easy to profile the pool

11/17 Reserve a worker for a task

```
func main() {  
    jobs := make(chan chan int) // first chan to reserve a worker, second one to send a task  
    for i := 0; i < 2; i++ {  
        go worker(jobs)  
    }  
  
    for i := 0; i < 10; i++ {  
        job := make(chan int)  
        select {  
            case jobs <- job:  
                fmt.Println("reserve worker for task:", i)  
                job <- i // build and send the task  
            default:  
                fmt.Println("drop task:", i)  
        }  
    }  
}  
  
func worker(jobs chan chan int) {  
    for job := range jobs {  
        task := <- job  
        fmt.Println("done task:", task)  
    }  
}
```

12/17 Prefer non-blocking sending

```
select {  
  case ch1 <- "value1": // strategy 1  
  case ch2 <- "value2": // strategy 2  
  default:  
    // take an action  
}
```

- No `time.After` to control the pool
- Use buffered channels for backpressure

13/17 Requirements for the pool

- Worker must know how to restart itself
- Fixed size of the pool and adjust per instance type
- Send metrics to centralized location and keep the current state on the instance

What to profile:

- How long does it take to process the task
- How many workers are busy right now
- What was the maximum number of concurrently occupied workers during lifetime of the process

14/17 Single buffer

```
type A struct { B int; _ int }

func main() {
    profile(func(){
        for i := 0; i < 10; i++ {
            list := make([]*A, 1000)
            for i := 0; i < 1000; i++ {
                list[i] = &A{B: i}
            }
        }
    })
}

func profile(fn func()){
    o := new(runtime.MemStats)
    runtime.ReadMemStats(o)
    fn()
    n := new(runtime.MemStats)
    runtime.ReadMemStats(n)
    fmt.Printf("objects %v alloc %v", n.HeapObjects-o.HeapObjects, n.HeapAlloc - o.HeapAlloc)
}
```

[Run](#)

15/17 Double buffers


```
func main() {  
    profile(func(){  
        lists := make([][]*A, 2)  
        lists[0] = make([]*A, 1000)  
        lists[1] = make([]*A, 1000)  
  
        for x := 0; x < 5; x++ {  
            for y := 0; y < 1000; y++ {  
                a := lists[0][y]  
                if a == nil {  
                    a = &A{ }  
                }  
                a.B = x*y  
                lists[0][y] = a  
            }  
  
            lists[0], lists[1] = lists[1], lists[0]  
        }  
    })  
}
```

Run

16/17 Sometimes we have to write our library

[pubnative/mysqldriver-go](https://github.com/pubnative/mysqldriver-go) (<https://github.com/pubnative/mysqldriver-go>) GC optimized MySQL driver

[pubnative/mysqlproto-go](https://github.com/pubnative/mysqlproto-go) (<https://github.com/pubnative/mysqlproto-go>) Heap friendly implementation of the MySQL protocol



17/17 Keep the test code in one function

```
func TestMyFunc(t *testing.T) {  
    // 1. setup test  
    // 2. define and execute test cases  
    // 3. cleanup  
}
```

- Keep tests straightforward, reading from top to bottom you should be able to understand how the code performs and what are requirements for it
- No helpers and abstractions in the test code
- Define what's required to execute the code you test, nothing more

Thank you

Kostiantyn Stepaniuk

PubNative GmbH

kostiantyn@pubnative.net (mailto:kostiantyn@pubnative.net)